

```

In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler

from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, classification_report, roc_auc_score, roc_curve
)

#Load Dataset
df =pd.read_csv('C:/Users/sarve/OneDrive/Desktop/Logistick Regression Healthcare Pr

#Preprocessing
df['ScheduledDay'] = pd.to_datetime(df['ScheduledDay'])
df['AppointmentDay'] = pd.to_datetime(df['AppointmentDay'])
df['WaitTime'] = (df['AppointmentDay'] - df['ScheduledDay']).dt.days
df['No-show'] = df['No-show'].map({'No': 0, 'Yes': 1})

# data cleaning drop irrelevant columns

df = df.drop(columns = ['PatientId', 'AppointmentID'])

# Remove negative wait times
df = df[df['WaitTime'] >= 0]

# Encode categorical variables
df = pd.get_dummies(df, columns=['Gender', 'Neighbourhood'], drop_first=True)

```

```

In [2]: df.head()

```

Out[2]:

	ScheduledDay	AppointmentDay	Age	Scholarship	Hipertension	Diabetes	Alcoholism
5	2016-04-27 08:36:51+00:00	2016-04-29 00:00:00+00:00	76	0	1	0	C
6	2016-04-27 15:05:12+00:00	2016-04-29 00:00:00+00:00	23	0	0	0	C
7	2016-04-27 15:39:58+00:00	2016-04-29 00:00:00+00:00	39	0	0	0	C
9	2016-04-27 12:48:25+00:00	2016-04-29 00:00:00+00:00	19	0	0	0	C
10	2016-04-27 14:58:11+00:00	2016-04-29 00:00:00+00:00	30	0	0	0	C

5 rows × 91 columns



```
In [3]: #Feature and Target
X= df.drop(columns =['No-show', 'ScheduledDay', 'AppointmentDay'])
y = df['No-show']
```

```
In [4]: # Train and Test Split data

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2, random_sta
```

```
In [5]: #Feature scaling
scaler =StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.fit_transform(X_test)
```

```
In [6]: #Model Training
model = LogisticRegression(max_iter=1000)
model.fit(X_train_scaled,y_train)
```

```
Out[6]: LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [ ]: #Predictions

y_pred = model.predict(X_test_scaled)

y_prob = model.predict_proba(X_test_scaled)[: ,1]
```

## Evaluating model performance

```
print("Accuracy:", round(accuracy_score(y_test, y_pred), 4)) print("Precision:",
round(precision_score(y_test,y_pred),4)) print("Recall:", round(recall_score(y_test, y_pred), 4))
print("F1 Score:", round(f1_score(y_test, y_pred), 4)) print("ROC AUC:",
round(roc_auc_score(y_test, y_prob), 4))
```

```
In [8]: # Improving class imbalance with with class_weights
model =LogisticRegression(max_iter=1000,class_weight='balanced')
model.fit(X_train_scaled,y_train)
```

```
Out[8]: LogisticRegression
LogisticRegression(class_weight='balanced', max_iter=1000)
```

```
In [9]: #Threshold Tuning
threshold = 0.2 # example threshold, tune via validation
y_pred_custom = (y_prob >= threshold).astype(int)

print("Recall at threshold 0.2:", round(recall_score(y_test, y_pred_custom), 4))
print("Precision at threshold 0.2:", round(precision_score(y_test, y_pred_custom),
print("F1 Score at threshold 0.2:", round(f1_score(y_test, y_pred_custom), 4))
```

```
Recall at threshold 0.2: 0.9276
Precision at threshold 0.2: 0.2946
F1 Score at threshold 0.2: 0.4472
```

```
In [10]: ## Over Sampling the minority class using SMOTE
from imblearn.over_sampling import SMOTE
```

```
In [11]: #using smote to resample the minority classes using synthetic resampling
sm =SMOTE(random_state=42)

#fit_resample function will perform synthetic resampling

X_train_res, y_train_res = sm.fit_resample(X_train_scaled, y_train)
```

```
In [12]: ### Training model again after performing the SMOTE
model =LogisticRegression(max_iter=1000)
model.fit(X_train_res,y_train_res)
```

```
Out[12]: LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [13]: ## model prediction
y_pred = model.predict(X_test_scaled)
y_proba = model.predict_proba(X_test_scaled)[: , 1]
```

```
In [14]: ## Evaluating Model predictions

print("Recall after SMOTE:", round(recall_score(y_test, y_pred), 4))
```

```
print("Precision after SMOTE:", round(precision_score(y_test, y_pred), 4))
print("F1 Score after SMOTE:", round(f1_score(y_test, y_pred), 4))
print("ROC AUC:", round(roc_auc_score(y_test, y_prob), 4))
```

Recall after SMOTE: 0.5691


Precision after SMOTE: 0.3421

F1 Score after SMOTE: 0.4273

ROC AUC: 0.5888

In [15]:

```
## Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



In [16]:

```
# Cross Validation to measure how well model would perform on unseen data
cv_scores = cross_val_score(model,X, y, cv=5, scoring='roc_auc')
print("Cross-validated ROC AUC scores:", np.round(cv_scores, 4))
```

Cross-validated ROC AUC scores: [0.5822 0.5789 0.5669 0.5611 0.5857]

In [17]:

```
print("Cross-validated ROC AUC scores:", np.round(cv_scores, 4))
```

Cross-validated ROC AUC scores: [0.5822 0.5789 0.5669 0.5611 0.5857]

In [18]:

```
print("Average CV ROC AUC:", round(np.mean(cv_scores), 4))
```

Average CV ROC AUC: 0.575

```
In [19]: patient_results = X_test.copy()
```

```
In [20]: patient_results["Predicted_Probability"] = y_prob
```

```
In [21]: patient_results
```

```
Out[21]:
```

	Age	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap	SMS_received	Wa
<b>734</b>	7	0	0	0	0	0	1	
<b>79659</b>	77	0	0	0	0	0	1	
<b>60687</b>	34	0	0	0	0	0	0	
<b>10620</b>	41	0	0	0	0	0	0	
<b>41582</b>	61	0	0	0	0	0	0	
...	...	...	...	...	...	...	...	...
<b>38562</b>	7	0	0	0	0	0	1	
<b>27660</b>	55	0	1	1	0	0	1	
<b>46560</b>	76	0	1	0	0	0	1	
<b>93780</b>	63	0	1	0	1	0	0	
<b>3200</b>	59	0	1	0	0	0	1	

14392 rows × 89 columns



```
In [22]: y_preds = model.predict(X_test) # 0 or 1
patient_results["Predicted_Label"] = y_preds
```

c:\Users\sarve\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2732: UserWarning: X has feature names, but LogisticRegression was fitted without feature names  
warnings.warn(

```
In [23]: patient_results
```

Out[23]:

	Age	Scholarship	Hipertension	Diabetes	Alcoholism	Handcap	SMS_received	Wa
734	7	0	0	0	0	0	1	
79659	77	0	0	0	0	0	1	
60687	34	0	0	0	0	0	0	
10620	41	0	0	0	0	0	0	
41582	61	0	0	0	0	0	0	
...	...	...	...	...	...	...	...	...
38562	7	0	0	0	0	0	1	
27660	55	0	1	1	0	0	1	
46560	76	0	1	0	0	0	1	
93780	63	0	1	0	1	0	0	
3200	59	0	1	0	0	0	1	

14392 rows × 90 columns



In [ ]: